

INTRODUCTION

A modern computer is a complex machine made up of various components that work together to perform a wide range of tasks. At the heart of every computer is one or more processors, also known as central processing units (CPUs), which are responsible for executing instructions and performing calculations.

In order to function, a computer needs a place to store data and instructions. This is where main memory, or RAM, comes in. Main memory is a type of temporary storage that holds data and instructions while they are being used by the processor.

A computer also needs a way to store data and programs permanently, even when the power is turned off. This is where disks, such as hard drives and solid-state drives, come in. These devices use magnetic or electronic storage to hold large amounts of data, allowing a computer to access and use it even when it is not turned on.

In addition to these core components, a modern computer also includes a variety of input and output devices. These include printers, which allow users to create physical copies of documents and other materials, keyboards and mice, which provide a way for users to input data and control the computer, and displays, which show the output of the computer.

Networks interfaces, such as Ethernet ports and wireless adapters, enable a computer to connect to the internet and other devices, while other input/output devices, such as scanners and webcams, provide additional functionality.

An operating system (OS) is a vital component of every computer. It is a collection of software that manages the hardware and software resources of a computer, and provides a platform for other programs to run on. Most readers will be familiar with popular operating systems such as Windows, Linux, FreeBSD, or OS X.

However, it is important to note that the program that users interact with, commonly referred to as the shell when it is text based and the GUI (Graphical User Interface) when it uses icons, is not actually part of the operating system.

The shell is a user interface that allows users to interact with the operating system by entering commands through a command line or terminal window. The GUI, on the other hand, is a more visual interface that allows users to interact with the operating system using icons, buttons, and other graphical elements.

While the shell and GUI are not part of the operating system itself, they do rely on the operating system to function. They use the resources and services provided by the operating system to carry out tasks and provide a user-friendly interface for interacting with the computer.

In a computer, hardware refers to the physical components that make up the machine. These include chips, boards, disks, a keyboard, a monitor, and other physical objects. Hardware is responsible for performing the basic functions of a computer, such as storing data, processing instructions, and displaying output.

On top of the hardware is the software, which consists of the programs and operating system that run on a computer. Software is responsible for controlling the hardware and performing tasks according to a set of instructions.

Most computers have two modes of operation: kernel mode and user mode. Kernel mode, also known as privileged mode, is a high-level mode of operation in which the operating system has complete control over the hardware and software resources of the computer. Kernel mode is used to perform critical tasks, such as managing memory and input/output operations.

User mode, on the other hand, is a lower-level mode of operation in which the operating system allows programs to access certain hardware and software resources. User mode is used to run non-critical tasks, such as running applications and programs.

The user interface is the program that allows users to interact with a computer and its operating system. There are two main types of user interfaces: the shell, which is a text-based interface, and the GUI (Graphical User Interface), which uses icons and other visual elements.

The shell is the lowest level of user-mode software, and allows users to enter commands through a command line or terminal window. The GUI, on the other hand, provides a more visual interface that allows users to interact with the operating system using icons, buttons, and other graphical elements.

Both the shell and GUI allow users to start other programs, such as a web browser, email reader, or music player. They provide a way for users to access and use the various software applications and tools installed on their computer.

In addition to starting programs, the user interface also allows users to perform a wide range of tasks, such as creating and editing documents, browsing the internet, and organizing files and folders. It is an essential part of the computing experience, providing a convenient and user-friendly way to interact with the computer.

An operating system is a vital component of every computer, responsible for managing the hardware and software resources of the machine. It provides a platform for other programs to run on, and serves as the interface between the hardware and the user.

Normal software, on the other hand, is any program that is not part of the operating system. This includes applications and tools such as email readers, web browsers, and productivity software.

An important distinction between the operating system and normal software is that users are free to choose which normal software they want to use, or even write their own if they have the necessary skills and resources. This is because normal software is not critical to the functioning of the operating system, and can be easily replaced or modified.

However, the operating system includes certain components, such as the clock interrupt handler, that are protected by hardware and cannot be modified by users. These components are critical to the functioning of the operating system, and any attempts to modify them could potentially cause serious problems.

In most computers, there is a clear distinction between the operating system and normal software. The operating system is a collection of software that manages the hardware and software resources of a computer, and provides a platform for other programs to run on. Normal software, on the other hand, is any program that is not part of the operating system, and can include applications and tools such as email readers, web browsers, and productivity software.

However, this distinction can become blurred in certain types of systems, such as embedded systems and interpreted systems.

Embedded systems are small, specialized computers that are used to perform specific tasks in devices such as smartphones, televisions, and appliances. These systems often do not have a traditional operating system with a kernel mode, and instead rely on specialized software to control the hardware and perform tasks.

Interpreted systems, such as those based on Java, use interpretation, rather than hardware, to separate the components of the system. In these systems, a program called an interpreter is used to execute instructions, rather than the hardware itself. This can make it more difficult to distinguish between the operating system and normal software, as the interpreter is responsible for executing both types of programs.

Operating systems are essential components of every computer, responsible for managing the hardware and software resources of the machine and providing a platform for other programs to run on. However, writing an operating system is a complex and time-consuming task, and it is not uncommon for an operating system to take years or even decades to develop.

There are several reasons why operating systems live a long time. First and foremost, they are very hard to write. Developing an operating system requires a deep understanding of computer hardware and software, as well as the ability to design and implement complex algorithms and systems. It is a challenging and specialized field, and writing an operating system is a major accomplishment.

In addition, once an operating system is written, the owner is often loath to throw it out and start again. This is because an operating system is the foundation of a computer, and changing it can be a risky and time-consuming process. It requires extensive testing and debugging to ensure that the new operating system is stable and compatible with the hardware and software of the computer.

Furthermore, once an operating system is in widespread use, it becomes deeply integrated into the ecosystem of software and hardware that relies on it. Changing the operating system would require extensive updates and modifications to these other components, which can be a costly and complex process.

What is Operating System

The Operating System as an Extended Machine

The architecture of a computer refers to the overall design and organization of the hardware and software components of the machine. It includes the instruction set, which is the set of instructions that the processor is capable of executing, as well as the memory organization, input/output (I/O) system, and bus structure.

At the machine-language level, the architecture of most computers is primitive and awkward to program, especially for input/output tasks. This is because machine language is the lowest level of programming, and consists of a series of binary instructions that are difficult for humans to read and understand. As a result, programming at the machine-language level is time-consuming and error-prone, and requires a deep understanding of the underlying hardware.

To make this point more concrete, consider modern SATA (Serial ATA) hard disks, which are used in most computers. These devices use a complex interface to transfer data between the hard disk and the computer, and programming this interface at the machine-language level would be a challenging and time-consuming task.

However, most users do not have to program at the machine-language level, thanks to higher-level programming languages and tools that provide a more user-friendly interface for developing software. These tools allow programmers to focus on the logic and functionality of their programs, rather than the details of the hardware and machine language.

Abstraction is a fundamental concept in computer science, and refers to the process of hiding the details of a system or process in order to focus on the essential aspects. Abstraction is used throughout computing to make complex systems more manageable and easier to understand.

One example of abstraction in computing is the file system, which is a way of organizing and storing data on a computer. The file system provides an interface for programs to create, write, and read files, without having to deal with the messy details of how the hardware actually works.

At the lowest level, a hard disk is simply a collection of magnetic or electronic storage media, and reading and writing data to it requires precise control of the hardware. However, the file system abstracts this complexity, providing a logical structure for organizing and accessing data that is easy for programs to use.

Using this abstraction, a program can create a file, write data to it, and read the data back, without having to worry about the physical location of the data on the hard disk or the details of how the hardware reads and writes data. This makes it much easier for programs to access and manipulate data, and enables the development of more complex and sophisticated software.

Abstraction is a fundamental concept in computer science, and is used to hide the complexity of systems and processes in order to focus on the essential aspects. This is especially important in modern computing, where systems and processes can be extremely complex and difficult to manage.

One of the key benefits of abstraction is that it turns a nearly impossible task into two manageable ones. The first is defining and implementing the abstractions. This involves identifying the essential aspects of a system or process, and creating an interface or representation that hides the unnecessary details. This can be a challenging task, as it requires a deep understanding of the underlying system and the ability to design and implement the abstraction in a way that is easy to use and understand.

The second task is using the abstractions to solve problems and perform tasks. Once the abstractions have been defined and implemented, they can be used to simplify and manage the complexity of the underlying system. This makes it much easier for programmers to develop software and perform tasks, and enables the development of more complex and sophisticated systems.

Hardware is a crucial part of every computer, and is responsible for performing the basic functions of the machine, such as storing data, processing instructions, and displaying output. However, despite the efforts of industrial engineers to design hardware in a logical and user-friendly way, it is often difficult and awkward to program.

There are several reasons why hardware can be difficult to program. First and foremost, real processors, memories, disks, and other devices are very complicated, with intricate and specialized components that can be difficult to understand and work with. This complexity is necessary to enable the hardware to perform its functions, but it can also make it challenging to write software that is optimized for the hardware.

Another reason why hardware can be difficult to program is that it often presents awkward, idiosyncratic, and inconsistent interfaces to software developers. This can be due to the need for backward compatibility with older hardware, which can require software developers to work around outdated or deprecated features. It can also be due to cost-saving measures, which can lead to hardware that is less user-friendly and more difficult to work with.

Overall, it is important to recognize that hardware is often ugly, in the sense that it is difficult and awkward to program. Despite the best efforts of industrial engineers, real processors, memories, disks, and other devices are very complicated and present challenging interfaces to software developers. Understanding this can help users appreciate the complexity of hardware, and the effort that goes into developing software that is optimized for it.

The operating system is a critical component of every computer, responsible for managing the hardware and software resources of the machine and providing a platform for other programs to run on. However, the operating system's real customers are the application programs, which are the programs that users interact with to perform tasks such as word processing, internet browsing, and email management. These programs rely on the abstractions provided by the operating system to access and use the hardware and software resources of the computer.

In contrast, end users, who are the people who use the computer, do not interact directly with the operating system and its abstractions. Instead, they use the abstractions provided by the user interface, which is the program that allows users to interact with the computer and its operating system. There are

two main types of user interfaces: the command-line shell, which is a text-based interface, and the graphical interface (GUI), which uses icons and other visual elements.

While the abstractions at the user interface may be similar to the ones provided by the operating system, this is not always the case. For example, a user may see a very different interface depending on whether they are using a normal desktop environment or a line-oriented command prompt, but the underlying operating system abstractions are the same in both cases. Similarly, a Linux user running Gnome or KDE may see a different interface than a Linux user working directly on top of the underlying X Window System, but the underlying operating system abstractions are the same in both cases.

The Operating System as a Resource Manager

There are two main ways to view the role of an operating system in a computer: a top-down view and a bottom-up view. In the top-down view, the operating system is primarily responsible for providing abstractions to application programs, which are the programs that users interact with to perform tasks such as word processing, internet browsing, and email management. This view emphasizes the role of the operating system in providing a platform for other programs to run on, and in hiding the complexity of the underlying hardware from the application programs.

In contrast, the bottom-up view holds that the operating system is primarily responsible for managing all the pieces of a complex system, such as processors, memories, timers, disks, mice, network interfaces, printers, and other devices. In this view, the job of the operating system is to provide for an orderly and controlled allocation of these resources among the various programs that want to use them. This view emphasizes the role of the operating system in managing the hardware and software resources of the computer, and in ensuring that these resources are used efficiently and effectively.

Both the top-down and bottom-up views of the operating system are valid, and reflect different aspects of the role of the operating system in a computer. The top-down view emphasizes the importance of abstractions in making it easier for application programs to use the resources of the computer, while the bottom-up view emphasizes the importance of resource management in ensuring the efficient and effective use of these resources. Understanding both views can help users appreciate the complexity of the operating system and the various roles it plays in managing the hardware and software resources of a computer.

Modern operating systems allow multiple programs to be in memory and run at the same time, which is a significant improvement over older operating systems that could only run one program at a time. This is made possible by the use of multiprogramming, which is a technique that allows the operating system to divide the processor's time among multiple programs.

However, running multiple programs simultaneously can also create problems if they try to access the same resource at the same time. For example, imagine what would happen if three programs running on some computer all tried to print their output simultaneously on the same printer. Without proper management,

the printer would likely become congested and the programs would have to wait their turn to print, leading to delays and inefficiencies.

To prevent this kind of resource contention, modern operating systems use a technique called resource scheduling, which is responsible for allocating the available resources of the computer to the programs that need them. This can include allocating processor time, memory, I/O devices, and other resources in a way that ensures that each program gets the resources it needs to function properly.

When a computer or network has more than one user, the need for managing and protecting the memory, I/O devices, and other resources becomes even more important, since the users might otherwise interfere with one another. Without proper management and protection, users could accidentally or intentionally access or modify the resources of the computer or network, leading to problems such as data loss, security breaches, and system instability.

To prevent these kinds of issues, modern operating systems use a variety of techniques to manage and protect the resources of the computer or network. These techniques can include resource scheduling, which is responsible for allocating the available resources of the computer to the programs that need them, and resource protection, which is responsible for restricting access to certain resources to authorized users only.

Resource scheduling and protection are especially important in multi-user environments, where there may be many users trying to access and use the resources of the computer or network simultaneously. By managing and protecting these resources, the operating system can ensure that each user gets the resources they need to perform their tasks, while also preventing interference and conflicts between users.

Resource management is a critical aspect of modern operating systems, and is responsible for allocating the available resources of the computer or network to the programs that need them. One of the key techniques used in resource management is multiplexing, which refers to the process of sharing resources in two different ways: in time and in space.

Time multiplexing is a technique used in resource management to share a resource among multiple programs or users by allocating the resource to one program or user at a time in a round-robin fashion. This is an effective way of maximizing the use of the resource, as it allows multiple programs or users to share it without having to wait for it to become available.

To understand how time multiplexing works, consider the example of a processor, which is a resource that can be shared among multiple programs. When a processor is time multiplexed, different programs take turns using it. First, one program gets to use the processor, then another, and so on. This allows each program to run for a short period of time, known as a time slice, before yielding to the next program.

One of the key benefits of time multiplexing is that it allows multiple programs to run concurrently, even if there is only one processor available. This can significantly improve the performance of the computer, as it allows the processor to be utilized more efficiently and reduces the time that programs have to wait to be executed.

Space multiplexing is a technique used in resource management to share a resource among multiple programs or users by dividing the resource into smaller units that can be allocated to different programs or users. This is an effective way of maximizing the use of the resource, as it allows multiple programs or users to share it without having to wait for it to become available.

To understand how space multiplexing works, consider the example of memory, which is a resource that can be shared among multiple programs. When memory is space multiplexed, each program gets a portion of the memory to use. This allows each program to run concurrently, even if there is not enough memory available to accommodate all of the programs at once.

One of the key benefits of space multiplexing is that it allows multiple programs to run concurrently, even if there is not enough memory available to accommodate all of the programs at once. This can significantly improve the performance of the computer, as it allows the memory to be utilized more efficiently and reduces the time that programs have to wait to be executed.

HISTORY OF OPERATING SYSTEMS

Operating systems are an essential component of every computer, responsible for managing the hardware and software resources of the machine and providing a platform for other programs to run on. Over the years, operating systems have undergone significant evolution, with new technologies and capabilities being added to meet the changing needs of users and the advancement of computer hardware. In this blog post, we will briefly look at a few of the highlights of this evolution.

One of the earliest operating systems was the UNIX operating system, which was developed in the 1970s at Bell Labs. UNIX was designed to be a multi-user, multitasking operating system, and was widely used in the academic and research communities. Over the years, UNIX has evolved into several different versions, such as Linux and BSD, which are still widely used today.

Another notable operating system is the Microsoft Windows operating system, which was first released in 1985. Windows was designed to be a graphical user interface (GUI) operating system, and was widely adopted due to its user-friendly interface and widespread availability. Over the years, Windows has undergone numerous updates and releases, with each new version adding new features and capabilities.

In recent years, there has been a trend towards mobile operating systems, such as Android and iOS, which are designed for use on smartphones and tablets. These operating systems are designed to be lightweight and efficient, and are optimized for use on small, portable devices.

The first true digital computer was designed by the English mathematician Charles Babbage in the 19th century. Babbage, who is considered the father of computers, spent most of his life and fortune trying to build his "analytical engine," which was a general-purpose mechanical computer that could perform a wide range of calculations.

Although Babbage's analytical engine was a revolutionary concept, it was never completed due to the limitations of the technology of the time. The analytical engine was purely mechanical, and required wheels, gears, and cogs to be built to a high precision in order to function properly. However, the manufacturing technology of the day was unable to produce these components to the required level of accuracy, and as a result, Babbage was never able to get his analytical engine working properly.

Despite this, Babbage's work laid the foundation for the development of modern computers, and his contributions to the field of computer science are still recognized today. His concept of the analytical engine, which was a programmable computer that could be used to perform a wide range of calculations, was ahead of its time and paved the way for the development of modern computers that are able to perform a wide range of tasks.

Charles Babbage's analytical engine, which was the first true digital computer, was a revolutionary concept that laid the foundation for the development of modern computers. In addition to designing the hardware for the analytical engine, Babbage also realized that he would need software in order to make the machine useful.

To this end, Babbage hired a young woman named Ada Lovelace as the world's first programmer. Lovelace was the daughter of the famed British poet Lord Byron, and was a talented mathematician in her own right. She worked with Babbage to develop algorithms and programs for the analytical engine, and is credited with writing the world's first computer program.

Lovelace's work with Babbage was groundbreaking, and she is now considered one of the pioneers of computer science. Her contributions to the field were far ahead of their time, and she is remembered today as the world's first programmer.

The First Generation (1945–55): Vacuum Tubes

After Charles Babbage's unsuccessful efforts to build the first digital computer in the 19th century, little progress was made in the field until the World War II period, which stimulated an explosion of activity in the field of computer science. One of the key figures in this period was Professor John Atanasoff, who worked with his graduate student Clifford Berry to build what is now regarded as the first functioning digital computer at Iowa State University.

Atanasoff and Berry's computer, which was known as the Atanasoff-Berry Computer (ABC), was a revolutionary machine that was designed to solve systems of simultaneous linear equations. The ABC was the first computer to use electronic switches and binary digits (bits) to represent data, and was also the first to use a separate memory unit to store data and instructions.

Although the ABC was never completed, it was a major milestone in the development of modern computers, and its concepts and technologies were later incorporated into other computers. Atanasoff and Berry's work is now recognized as an important precursor to the development of modern computers, and they are remembered as two of the pioneers of the field.

The period during and after World War II saw an explosion of activity in the field of computer science, with many key figures working to develop the first digital computers. At roughly the same time as the Atanasoff-Berry Computer (ABC) was being developed at Iowa State University, Konrad Zuse in Berlin was building the Z3 computer out of electromechanical relays.

Other notable computers developed during this period include the Colossus, which was built and programmed by a group of scientists (including Alan Turing) at Bletchley Park, England in 1944. The Colossus was the first programmable electronic computer, and was used to break the German Enigma code during the war.

Other early computers include the Mark I, which was built by Howard Aiken at Harvard University in 1944, and the ENIAC, which was built by William Mauchley and his graduate student J. Presper Eckert at the University of Pennsylvania in 1945. The ENIAC was the first general-purpose electronic computer, and was capable of being programmed to perform a wide range of tasks.

In the early days of computers, the field was much different than it is today. A single group of people, usually engineers, were responsible for designing, building, programming, operating, and maintaining each machine. This was a very labor-intensive process, and required a high level of expertise in a wide range of fields.

All programming in these early computers was done in absolute machine language, which required the programmer to write code directly in the language that the computer's hardware could understand. This was a very tedious and error-prone process, and required a high level of skill and knowledge on the part of the programmer.

Even worse than machine language was the process of wiring up electrical circuits by connecting thousands of cables to plugboards in order to control the machine's basic functions. This process was even more labor-intensive and error-prone than machine language programming, and required a high level of expertise in electrical engineering.

In these early days, programming languages were unknown, and even assembly language, which is a low-level programming language that is easier to use than machine language, was not yet developed. Operating systems, which are software programs that manage the resources of a computer and provide a platform for other programs to run on, were also unheard of.

By the early 1950s, the process of programming computers had improved somewhat with the introduction of punched cards. Punched cards were a simple and effective way of inputting data and instructions into a computer, and made it much easier to write and run programs.

Instead of using plugboards or manually wiring up electrical circuits to control the machine's basic functions, programmers could now write their programs on punched cards and read them into the computer using a card reader. This made the process of programming computers much easier and more

efficient, and allowed programmers to focus on the logic of their programs rather than the tedious details of how the computer worked.

Despite this improvement, the process of programming computers in the early 1950s was still very labor-intensive and required a high level of skill and knowledge on the part of the programmer. Programming languages and operating systems had not yet been developed, and all programming was still done in machine language or assembly language.

The Second Generation (1955–65): Transistors and Batch Systems

The introduction of the transistor in the mid-1950s was a major turning point in the history of computers, and changed the field radically. The transistor, which is a small electronic device that can amplify or switch electronic signals, made it possible to build computers that were smaller, faster, and more reliable than ever before.

As a result of the transistor, computers became reliable enough that they could be manufactured and sold to paying customers with the expectation that they would continue to function long enough to get some useful work done. This marked the beginning of the computer industry as we know it today, and made it possible for computers to be used in a wide range of applications.

The introduction of the transistor also marked the first time that there was a clear separation between the designers, builders, operators, programmers, and maintenance personnel of computers. This separation of roles allowed for a more specialized and efficient development process, and made it possible for computers to be used in a wider range of applications.

In the early days of the computer industry, computers were very large and expensive machines known as mainframes. Mainframes were locked away in large, specially air-conditioned computer rooms, and required staffs of professional operators to run them.

Due to their size and cost, mainframes were only affordable for large corporations, major government agencies, or universities. The multimillion-dollar price tag for a mainframe made it out of reach for most small businesses or individuals, and as a result, only a select few were able to use these powerful machines.

Mainframes were used to perform a wide range of tasks, including scientific research, data processing, and business applications. They were also used to run large-scale computer systems, such as those used for airline reservations or banking.

In the early days of computers, the process of running a program was a very labor-intensive process. When a computer finished running a job, an operator would go over to the printer and tear off the output, and carry it over to the output room so that the programmer could collect it later.

The operator would then take one of the card decks that had been brought from the input room and read it into the computer. If the FORTRAN compiler, which is a software program that translates a program written in FORTRAN into machine language, was needed, the operator would have to get it from a file cabinet and read it into the computer.

This process was very time-consuming, and much computer time was wasted while operators were walking around the machine room. The operator had to physically move the card decks and other materials from one location to another, which was a very slow and error-prone process.

In the early days of computers, the high cost of the equipment meant that people were always looking for ways to reduce wasted time and make the most of their investment. One solution that was widely adopted was the batch system, which was a way of collecting and processing a tray full of jobs at once.

The batch system worked by collecting a tray full of jobs in the input room and then reading them onto a magnetic tape using a small, relatively inexpensive computer such as the IBM 1401. The IBM 1401 was good at reading cards, copying tapes, and printing output, but was not well-suited for numerical calculations.

Once the jobs were read onto the magnetic tape, they could be run on the larger, more expensive mainframe computer. The mainframe would process the jobs and produce the output, which was then printed and collected by the programmer.

The batch system was a way of reducing the wasted time associated with running jobs on a computer, and made it possible to get more work done with the same amount of resources. It was a key innovation in the early days of computers, and helped to make the most of the limited resources that were available.

In the second generation of computers, large machines were used primarily for scientific and engineering calculations, such as solving the partial differential equations that often occur in physics and engineering. These computers were largely programmed in FORTRAN and assembly language, which are programming languages that are used for scientific and engineering applications.

Typical operating systems for second-generation computers included FMS (the Fortran Monitor System) and IBSYS, which was IBM's operating system for the 7094. These operating systems were designed to support the scientific and engineering calculations that these computers were used for, and provided a range of tools and services to make it easier to develop and run programs on these machines.

The Third Generation (1965–1980): ICs and Multiprogramming

In the early 1960s, most computer manufacturers had two distinct product lines that were incompatible with one another. These product lines included word-oriented, large-scale scientific computers and character-oriented, commercial computers.

Word-oriented, large-scale scientific computers were used for industrial-strength numerical calculations in science and engineering, and were designed to handle large amounts of data and perform complex calculations. Examples of these computers included the 7094, which was widely used in scientific and engineering applications.

Character-oriented, commercial computers, on the other hand, were designed for use in business applications such as tape sorting and printing. These computers were widely used by banks and insurance companies to process large amounts of data and produce reports and other documents. Examples of these computers included the 1401, which was widely used in business applications.

In the early 1960s, IBM attempted to solve two problems that were facing the computer industry at the time. The first problem was the lack of compatibility between different types of computers, which made it difficult to write and run programs on more than one machine. The second problem was the need to have different types of computers for different applications, which was inefficient and costly.

To solve these problems, IBM introduced the System/360, which was a series of software-compatible machines ranging in size and power from small 1401-sized models to larger, more powerful machines like the 7094. The System/360 machines differed only in price and performance, with differences in maximum memory, processor speed, and the number of I/O devices permitted.

Since the System/360 machines all had the same architecture and instruction set, programs written for one machine could run on all the others, at least in theory. This made it much easier to write and run programs on multiple machines, and allowed users to choose the machine that was best suited for their needs based on their budget and performance requirements.

The IBM 360 was the first major computer line to use small-scale integrated circuits (ICs), which provided a significant price and performance advantage over the second-generation machines that were built using individual transistors. Integrated circuits are small, complex circuits that are made using a single piece of semiconductor material and are used to perform a variety of different functions in electronic devices. The use of integrated circuits in the IBM 360 allowed the machine to be much smaller and more efficient than earlier computers, and made it possible to produce computers that were more affordable and had better performance than ever before. Overall, the use of integrated circuits in the IBM 360 was a major innovation that helped to drive the widespread adoption of computers in a variety of industries, and laid the foundation for the development of modern computing technology.

The greatest strength of the "single-family" concept used in the IBM 360 was also its greatest weakness. The idea behind the single-family concept was that all software, including the operating system (OS/360), had to be able to run on all models of the IBM 360. This meant that the operating system had to be able to run on small systems, which were often used to replace 1401s for copying cards to tape, as well as on very large systems, which were used to replace 7094s for tasks like weather forecasting and other heavy computing.

While this approach had some advantages, it also had some significant limitations. In particular, it meant that the operating system had to be designed to work on a wide range of different hardware configurations, which made it more complex and harder to maintain. Additionally, the fact that the operating system had to work on both small and large systems meant that it had to be able to handle a wide range of workloads, which made it less efficient and less optimized for specific tasks.

IBM's operating system, known as OS/360, was a massive and complex piece of software that was designed to meet the needs of a wide range of users and applications. When it was first released in the 1960s, it was hailed as a revolutionary new platform that would enable businesses and organizations to do more with computers than ever before.

However, the development of OS/360 was not without its challenges. One major issue was the fact that IBM had to meet a wide range of conflicting requirements from its customers. Some users needed a simple, easy-to-use system, while others required a more powerful and flexible platform. This meant that IBM had to try to balance the needs of different groups of users, which was no easy task.

In addition to this, the sheer size and complexity of OS/360 made it extremely difficult to develop. The system consisted of millions of lines of code, written by thousands of programmers, and it contained thousands upon thousands of bugs. These bugs had to be constantly identified and corrected, which required a continuous stream of new releases to be made available.

OS/360 and other third-generation operating systems were able to meet the needs of most of their customers, despite their size and complexity. One of the key features that made these systems so successful was the use of multiprogramming.

In previous, second-generation operating systems, when a job needed to wait for input/output (I/O) to be completed, the central processing unit (CPU) would simply sit idle until the I/O was finished. This was not a significant problem for scientific calculations, where I/O operations were infrequent, but it was a major issue for commercial data processing. In these cases, the time spent waiting for I/O could be as much as 80-90% of the total processing time, which meant that the CPU was idle for much of the time.

To address this problem, third-generation operating systems introduced the concept of multiprogramming. With multiprogramming, the operating system could switch between multiple jobs, allowing the CPU to stay busy even when one job was waiting for I/O to be completed. This made it possible to get more work done in less time, and helped to improve the efficiency of the system.

Overall, the introduction of multiprogramming and other key techniques was a major step forward for operating systems, and helped to make them more powerful and efficient. It is no wonder that these systems were so widely adopted and have had such a lasting impact on the world of computing.

Third-generation operating systems introduced several important features that made them more efficient and user-friendly than their predecessors. One of these features was the ability to read jobs from cards onto a disk as soon as they were brought to the computer room. This process, known as spooling

(Simultaneous Peripheral Operation On Line), allowed jobs to be loaded onto the system and processed as soon as they were ready, rather than having to wait for a dedicated computer to become available.

Spooling was also used for output, allowing users to print or save the results of their jobs as soon as they were completed. With spooling, it was no longer necessary to have dedicated computers (such as the 1401) to handle input and output, which made it possible to streamline the process of running jobs on the system.

Overall, the introduction of spooling was a major improvement for third-generation operating systems, and it helped to make them more efficient and easier to use. The ability to load jobs onto the system as soon as they were ready, and to output the results of those jobs immediately, made it possible to get more work done in less time, and helped to make the process of using a computer more efficient and convenient.

The desire for quick response times led to the development of timesharing, a variant of multiprogramming that allowed multiple users to access a computer simultaneously through online terminals. In a timesharing system, the central processing unit (CPU) is allocated to each user in turn, allowing them to issue commands and receive immediate feedback.

This is especially useful for users who are debugging programs, as they can issue short commands (such as "compile this five-page procedure") and receive a quick response, rather than having to wait for longer, batch-processing jobs to be completed. Timesharing systems also allow the CPU to work on multiple tasks at once, allowing it to provide fast, interactive service to many users while also working on larger batch jobs in the background.

The first general-purpose timesharing system, called CTSS (Compatible Time Sharing System), was developed at the Massachusetts Institute of Technology (MIT) on a modified 7094 computer. However, timesharing did not become widely popular until the necessary protection hardware was more widespread during the third generation of operating systems.

After the success of the CTSS timesharing system, a group consisting of researchers from the Massachusetts Institute of Technology (MIT), Bell Labs, and General Electric (a major computer manufacturer at the time) decided to embark on the development of a "computer utility" that would be able to support hundreds of simultaneous timesharing users. This system, known as MULTICS (MULTIplexed Information and Computing Service), was designed to be a huge machine that would provide computing power to everyone in the Boston area.

The goal of the MULTICS project was to create a system that would be able to handle the computing needs of a large and diverse group of users, ranging from researchers and scientists to business professionals and everyday individuals. To achieve this, the designers of MULTICS envisioned a system that would be highly reliable, scalable, and flexible, with the ability to support a wide range of applications and services.

Despite the ambitious goals of the MULTICS project, the system was eventually deemed a failure, as it was unable to meet the needs of its users and was unable to achieve the level of adoption that was originally hoped for. However, the ideas and concepts that were developed as part of the MULTICS project went on

to influence the development of many other operating systems and technologies, and it remains an important part of the history of computing.

MULTICS was a mixed success when it was first released in the 1970s. While it was designed to support hundreds of users on a machine that was only slightly more powerful than a modern-day Intel 386-based personal computer, it had much more input/output (I/O) capacity. This was not as unrealistic as it might seem, as in those days programmers were skilled at writing small, efficient programs, a skill that has since been largely lost.

There were many reasons that MULTICS did not achieve the level of adoption that was originally hoped for. One issue was the fact that it was written in the PL/I programming language, which had a number of problems and was not as popular as other languages at the time. In addition, the PL/I compiler was years late in arriving and barely worked at all when it was finally released.

Despite these issues, MULTICS did go on to influence the development of many other operating systems and technologies. It was a pioneering system that helped to pave the way for the development of more advanced and powerful operating systems in the future. While it may not have taken over the world as originally intended, it remains an important part of the history of computing.

By the end of the 20th century, the idea of a computer utility had largely faded away, as advances in technology and the proliferation of personal computers made it easier for individuals and organizations to have their own computing resources. However, the concept of a computer utility may be making a comeback in the form of cloud computing.

In cloud computing, relatively small devices such as smartphones, tablets, and laptops are connected to servers in remote data centers where all of the computing is done. These devices act as the user interface, allowing people to access the computing power of the data center from wherever they are. The motivation for this model is that many people do not want to deal with the complexities of managing their own computer systems, and would prefer to have this work done by professionals working for the company running the data center.

The third generation of computing also saw the emergence of minicomputers, which were smaller and more affordable than mainframe computers but still capable of handling a wide range of applications. The DEC PDP-1, which was released in 1961, was one of the first minicomputers and quickly became popular due to its relatively low price (around \$120,000) and good performance for certain types of non-numerical work.

The PDP-1 had only 4K of 18-bit words, which was a small amount of memory compared to mainframe computers of the time, but it was still able to perform many tasks almost as fast as the more expensive mainframes. The success of the PDP-1 led to the development of a series of other PDP minicomputers, including the PDP-11, which became widely used in a variety of settings.

UNIX is a popular operating system that was originally developed by Ken Thompson, a computer scientist at Bell Labs who had worked on the MULTICS project. Thompson wanted to create a stripped-down, single-user version of MULTICS that could be used on a small PDP-7 minicomputer that was not being used for anything else.

Over time, Thompson's work on UNIX grew into a full-fledged operating system that became popular in the academic world, as well as with government agencies and many companies. One of the key features of UNIX was its portability, which meant that it could be easily adapted to run on a wide range of hardware platforms. This made it an attractive option for users who wanted a stable and reliable operating system that could be used on a variety of computers.

UNIX has had a lasting impact on the world of computing, and many of the ideas and concepts that were developed as part of the UNIX project have been incorporated into other operating systems and technologies. Today, UNIX and its many variants (such as Linux) are still widely used and are an important part of the computing landscape.

UNIX has evolved over time, and there are now several major versions of the operating system in use. Two of the most popular versions are System V, which was developed by AT&T, and BSD (Berkeley Software Distribution), which was developed by the University of California at Berkeley. Both of these versions have minor variants as well.

To make it easier for developers to write programs that could run on any version of UNIX, the Institute of Electrical and Electronics Engineers (IEEE) created a standard called POSIX (Portable Operating System Interface) that defines a minimal system-call interface that must be supported by all conforming UNIX systems. This standard has helped to ensure that UNIX programs can be easily ported between different versions of the operating system, making it easier for developers to create software that can be used by a wide audience.

In addition to being supported by most versions of UNIX, the POSIX interface has also been adopted by some other operating systems, further increasing its reach and making it an important part of the computing landscape. Overall, the development of POSIX has helped to make it easier for developers to create software that can be used on a wide range of platforms, and has played a key role in the evolution of UNIX and other operating systems.

MINIX is a small clone of the UNIX operating system that was developed by Andrew S. Tanenbaum in 1987 for educational purposes. Functionally, MINIX is very similar to UNIX, including support for the POSIX standard.

Since its initial release, MINIX has evolved into MINIX 3, which is a highly modular operating system that is focused on achieving very high levels of reliability. MINIX 3 is designed to be easy to use and understand, making it an attractive option for educational settings and for users who are new to operating systems.

Overall, MINIX is an important part of the history of UNIX and has helped to introduce many people to the concepts and ideas behind operating systems. Its development has contributed to the evolution of UNIX

and has helped to make it easier for people to learn about and understand the underlying principles of computer systems.

The Fourth Generation (1980–Present): Personal Computers

The development of Large Scale Integration (LSI) circuits, which are chips containing thousands of transistors on a small area of silicon, marked the beginning of the age of the personal computer. Personal computers, also known as microcomputers, were similar in architecture to minicomputers such as the PDP-11, but they were much more affordable and accessible to individuals.

Prior to the development of LSI circuits, it was only possible for larger organizations or departments to afford their own computers. With the advent of microprocessors, however, it became possible for a single individual to have their own personal computer. This revolutionized the way people used computers and made it possible for a much wider range of individuals and organizations to take advantage of the power and capabilities of these machines.

In 1974, Intel released the 8080, the first general-purpose 8-bit central processing unit (CPU). As part of the development process, Intel wanted to create an operating system for the 8080 that could be used to test the performance of the CPU. To achieve this, Intel hired Gary Kildall, a consultant, to write an operating system for the 8080.

Kildall and a friend first created a controller for the newly released Shugart Associates 8-inch floppy disk and connected it to the 8080, creating the first microcomputer with a disk drive. Kildall then wrote a disk-based operating system called CP/M (Control Program for Microcomputers) for the 8080.

Intel did not believe that disk-based microcomputers had much of a future, and when Kildall asked for the rights to CP/M, Intel granted his request. Kildall then formed a company called Digital Research to further develop and sell CP/M. CP/M went on to become a popular operating system for microcomputers and played a key role in the early history of personal computing.

In 1977, Digital Research released a revised version of CP/M that was optimized for use on microcomputers with 8080, Zilog Z80, and other CPU chips. This version of CP/M was widely adopted and became the dominant operating system for microcomputers for about five years. Many application programs were written to run on CP/M, which helped to increase its popularity and solidify its position as the dominant operating system for microcomputers.

Overall, CP/M played a significant role in the early history of personal computing, and its success helped to establish Digital Research as a major player in the field of operating system development. The release of the revised version of CP/M in 1977 was a key milestone in the evolution of the operating system and helped to make it even more widely available and useful to users.

In the early 1980s, IBM developed the IBM PC and began looking for software to run on it. In the course of this process, IBM contacted Bill Gates to license his BASIC interpreter, and also asked him if he knew of an operating system that could be used on the PC. Gates suggested that IBM contact Digital Research, which was then the dominant operating system company in the world.

However, Gary Kildall, the owner of Digital Research, refused to meet with IBM and instead sent a subordinate to represent the company. Kildall's lawyer even refused to sign IBM's nondisclosure agreement covering the not-yet-announced PC. As a result of this decision, IBM went back to Gates and asked if he could provide them with an operating system. This led to the development of MS-DOS (Microsoft Disk Operating System), which became the dominant operating system for the IBM PC and helped to establish Microsoft as a major player in the world of personal computing.

When IBM approached Microsoft about finding an operating system for the IBM PC, Bill Gates realized that a local computer manufacturer called Seattle Computer Products had a suitable operating system called DOS (Disk Operating System). He approached Seattle Computer Products and reportedly bought the operating system for \$75,000.

Gates then offered IBM a package that included both DOS and a BASIC interpreter, which IBM accepted. IBM requested some modifications to the operating system, so Gates hired the person who had written DOS, Tim Paterson, to make the changes. The revised system was renamed MS-DOS (MicroSoft Disk Operating System) and quickly became the dominant operating system for the IBM PC.

The success of MS-DOS helped to establish Microsoft as a major player in the world of personal computing, and the company has remained a dominant force in the industry to this day. The development of MS-DOS was a key moment in the evolution of personal computers and had a significant impact on the way that people use computers today.

By the time the IBM PC/AT, the successor to the IBM PC, was released in 1983 with the Intel 80286 CPU, MS-DOS had already established itself as the dominant operating system for personal computers. MS-DOS was later used on the 80386 and 80486 CPU chips as well. Although the initial version of MS-DOS was relatively basic, subsequent versions included more advanced features that were inspired by UNIX, a popular operating system developed at Bell Labs.

Microsoft had a long history with UNIX, and even sold a microcomputer version of the operating system called XENIX during the company's early years. The inclusion of UNIX-inspired features in later versions of MS-DOS helped to make it a more powerful and versatile operating system, and contributed to its widespread adoption.

CP/M, MS-DOS, and other operating systems for early microcomputers all used a command-line interface, which required users to type in commands using the keyboard. This changed due to the research of Doug Engelbart at Stanford Research Institute in the 1960s, who developed the concept of the graphical user interface (GUI).

A GUI is a type of operating system interface that uses visual elements such as windows, icons, and menus to allow users to interact with the computer. It also typically includes a pointing device such as a mouse, which can be used to select and manipulate on-screen elements. These ideas were later adopted by researchers at Xerox PARC and incorporated into the computers they developed.

The development of the GUI was a major milestone in the history of operating systems, as it made it much easier for users to interact with computers and greatly increased the usability of these systems. The GUI has since become the dominant type of operating system interface, and is used on almost all modern computers and devices.

One day, Steve Jobs, co-founder of Apple Inc., visited Xerox PARC and saw a demonstration of their graphical user interface (GUI). He immediately recognized the potential value of the GUI, which Xerox management had famously failed to see. This mistake was later chronicled in a book called "Fumbling the Future" (Smith and Alexander, 1988).

Inspired by the GUI, Jobs set out to build an Apple computer with a GUI of its own. This project led to the development of the Lisa, a computer that was too expensive and ultimately failed commercially. However, Jobs did not give up and continued to work on developing a GUI-based computer that would be more successful. This effort eventually resulted in the Apple Macintosh, which became a huge success due in part to its low price and user-friendly interface, which made it accessible to users who had no prior knowledge of computers or no desire to learn about them.

The development of the Macintosh was a major milestone in the evolution of personal computing and helped to establish Apple as a leading player in the industry. Its user-friendly interface and low price made it an appealing choice for many users, and contributed to its widespread adoption.

When Microsoft set out to build a successor to MS-DOS, it was heavily influenced by the success of the Apple Macintosh and its graphical user interface (GUI). As a result, Microsoft developed a GUI-based system called Windows, which originally ran on top of MS-DOS (i.e., as a shell rather than a true operating system).

From 1985 to 1995, Windows was primarily a graphical environment that ran on top of MS-DOS. However, starting in 1995, a standalone version of Windows called Windows 95 was released. This version incorporated many operating system features into it, and used MS-DOS only for booting and running older MS-DOS programs. In 1998, a slightly modified version of this system called Windows 98 was released. Both Windows 95 and Windows 98 still contained a significant amount of 16-bit Intel assembly language.

Overall, the development of Windows helped to popularize the GUI as the dominant type of operating system interface, and its success played a major role in the evolution of personal computing. The introduction of standalone versions of Windows, such as Windows 95 and 98, helped to make it a more powerful and feature-rich operating system, and contributed to its widespread adoption.

In 2001, Microsoft released Windows Me (Millennium Edition), which was a slightly upgraded version of Windows 98. This version of Windows included some new features and improvements over its predecessor, such as a system restore function and support for USB 2.0 devices. However, it was not as well received as earlier versions of Windows and was criticized for being prone to crashes and other issues.

In the same year, Microsoft also released Windows 2000, a version of Windows that was designed for business use. Windows 2000 was based on the same core technology as Windows 98 and Me, but it included a number of additional features and improvements, such as support for active directory and networking capabilities. It was generally well received and saw widespread adoption in the business world.

Overall, the release of Windows Me and Windows 2000 marked a continuation of Microsoft's efforts to improve and evolve the Windows operating system. These versions introduced new features and capabilities that helped to make them more powerful and useful for users.

After the release of Windows 2000, Microsoft continued to evolve the Windows operating system and introduced a number of new versions. The company divided the Windows family into two main lines: a client line, which was based on Windows XP and its successors, and a server line, which included Windows Server 2003 and Windows Server 2008. A third line for the embedded world was also released at a later date.

Each of these versions of Windows spawned numerous variations in the form of service packs, which introduced new features and improvements. This rapid release of new versions and service packs can be overwhelming for some administrators and those who write about operating systems. It can be challenging to keep track of all the different versions and their various features and capabilities.

UNIX is a widely-used operating system that is particularly popular on network and enterprise servers. It is known for its stability, security, and versatility, which make it well-suited for use in these types of environments. UNIX is also often used on desktop computers, notebooks, tablets, and smartphones, particularly in academic and research settings.

In recent years, Linux has emerged as a popular alternative to Windows on x86-based computers. Linux is a free and open-source operating system that is based on the UNIX kernel. It offers many of the same features and capabilities as UNIX, but it is typically more lightweight and less resource-intensive than its predecessor.

Linux has gained popularity among students and corporate users due to its low cost and wide range of available applications. It is also known for its strong security features and the ability to customize and modify the operating system to fit specific needs. As a result, Linux has become a viable alternative to Windows for many users.

FreeBSD is a popular open-source operating system that is derived from the BSD (Berkeley Software Distribution) project at the University of California, Berkeley. It is known for its stability, performance, and security, which make it well-suited for use on servers, workstations, and other specialized systems.

FreeBSD is widely used on workstations powered by high-performance RISC (Reduced Instruction Set Computing) chips. It is also the foundation for the OS X operating system used on modern Macintosh computers.

UNIX and its derivatives, including FreeBSD and Linux, are also used on a wide range of mobile devices, such as smartphones and tablets. For example, the iOS operating system used on Apple devices is based on the FreeBSD kernel, while the Android operating system used on many devices is based on the Linux kernel.

Network operating systems are designed to allow multiple computers to connect and communicate with each other over a network during the mid-1980s. With a network operating system, users can log in to remote machines and access resources, such as files and printers, as if they were local to their own machine.

Distributed operating systems, on the other hand, go a step further by allowing multiple computers to work together as a single system. In a distributed operating system, resources, such as processing power and storage, can be shared and accessed by all computers on the network. This allows for more efficient use of resources and can improve the overall performance of the system.

The growth of networks of personal computers running network and distributed operating systems has had a significant impact on the way we use and interact with computers. It has made it possible for people to access and share resources from anywhere, at any time, and has contributed to the rise of the internet and the proliferation of connected devices. Overall, the use of network and distributed operating systems has greatly increased the connectivity and collaboration possibilities for both individuals and organizations.

The Fifth Generation (1990–Present): Mobile Computers

A smartphone is a mobile device that combines the capabilities of a computer with the functionality of a phone. The first smartphone, the Nokia N9000, was released in the mid-1990s and featured a phone and a personal digital assistant (PDA) in one device. The term "smartphone" was coined by Ericsson in 1997 with the release of its GS88 "Penelope" device.

Since the introduction of the first smartphone, the technology has evolved significantly. Today's smartphones are much more advanced and powerful than their early counterparts, with a range of features and capabilities that go beyond basic phone and PDA functions.

Smartphones are equipped with a variety of sensors, including GPS, accelerometers, and cameras, and are capable of running a wide range of apps and accessing the internet. They have become an integral part of our daily lives, providing us with instant access to information, communication, and entertainment.

The smartphone market has become highly competitive, with many different manufacturers offering a wide range of devices to meet the needs and preferences of different users. Some of the most popular smartphone brands include Apple, Samsung, Huawei, and Xiaomi.

In 2010, Symbian's market share began to decline rapidly as Android and iOS gained popularity. By 2014, Symbian was no longer a major player in the smartphone market. Blackberry, too, has seen its market share decline in recent years, with Android and iOS now dominating the market.

Despite the dominance of Android and iOS, there are still many other operating systems in use on smartphones today. Microsoft's Windows Phone, while not as popular as the other two, still has a significant presence in certain markets. There are also smaller players like Tizen (developed by Samsung) and Sailfish (developed by Jolla).

The history of operating systems has been marked by the constant evolution of technology and the desire to improve user experience and efficiency. From the early days of batch processing and multiprogramming, to the development of GUI-based systems and the rise of smartphone operating systems, the field has seen a wealth of innovations. Each new generation of operating systems has brought with it new features and techniques that have shaped the way we use computers and other devices today. As technology continues to advance, it will be interesting to see what the future holds for operating systems and how they will continue to shape the way we interact with computers.

COMPUTER HARDWARE REVIEW

An operating system is a key component of any computer system, as it serves as the interface between the hardware and the software. It is responsible for managing the resources of the computer, such as the CPU, memory, and storage, and for providing a platform for applications to run on. The operating system extends the instruction set of the computer by providing additional functionality, such as system calls and libraries, that allow applications to access the hardware in a more convenient and abstracted way. As a result, the operating system must have a deep understanding of the hardware it is running on, including its capabilities, limitations, and the ways in which it appears to the programmer.

"An operating system (OS) is a set of software that manages the hardware and software resources of a computer. It is responsible for allocating these resources to the various programs and processes running on the computer. The OS provides an interface between the user and the computer, allowing the user to interact with the computer and its programs.

One of the main functions of an operating system is to manage the computer's memory. This involves allocating memory to different programs and processes, as well as managing virtual memory and swapping data to and from the hard drive as needed. The OS also manages the input and output (I/O) devices, such as the keyboard, mouse, and display, and coordinates the transfer of data between the CPU and these devices.

In addition to these basic functions, modern operating systems also provide a wide range of other features, such as support for networking, security, and user management. They may also include a range of utilities and tools for tasks such as file management, system maintenance, and debugging.

Processors

"The instructions that a CPU can execute are determined by its instruction set. The instruction set is the set of basic operations that the CPU can perform. Examples of instructions that might be found in a CPU's instruction set include adding two numbers together, moving a block of data from one location in memory to another, or comparing a value stored in memory to a value stored in a register.

In addition to its instruction set, the CPU also has a number of registers. Registers are small amounts of memory that are built into the CPU and are used to store data that the CPU is currently working on. For example, the CPU might use a register to store the operands for an instruction before performing the operation.

The CPU fetches instructions and data from memory and stores the results of its operations back in memory. Memory is made up of a large number of cells, each of which can store a single byte of data. The CPU can access any location in memory by specifying its address.

I/O devices are used to get data into and out of the computer. Examples of I/O devices include keyboards, mice, printers, and monitors. The CPU communicates with I/O devices using special instructions in its instruction set and through specialized registers known as I/O ports.

The operating system is responsible for managing the resources of the computer, including the CPU, memory, and I/O devices. It does this by controlling the execution of programs and allocating resources to them as needed. The operating system also provides a user interface, allowing users to interact with the computer and run programs."

The CPU is the central processing unit of the computer, responsible for executing instructions and carrying out the actions required by the operating system and other programs. Memory is used to store instructions and data for the CPU to access, while I/O devices are used to input and output data to and from the computer. The system bus connects these components and allows them to communicate with each other. Different CPU architectures, such as x86 and ARM, are not compatible with each other and require different operating systems and programs.

Another important register is the stack pointer, which points to the top of the stack in memory. The stack is used for temporary storage and for function calls. When a function is called, the arguments and the return address are pushed onto the stack. When the function returns, the return value is stored in a designated register and the stack is popped back to its previous state.

Other special registers may include the status register, which stores information about the current state of the CPU, and the memory management registers, which are used to implement virtual memory.

In addition to the registers and buses mentioned earlier, computers also have a variety of other hardware components, such as memory, I/O devices, and storage devices. Memory stores data and instructions that the CPU can access, while I/O devices allow the computer to communicate with the outside world. Storage devices, such as hard drives and SSDs, provide long-term storage for data and programs.

The operating system is responsible for managing these hardware resources and making them available to programs. It does this through a variety of system calls, which provide an interface between the hardware and the software. By using these system calls, programs can request access to hardware resources and the operating system can grant or deny these requests based on the current state of the system.

The stack is an area of memory that is used for storing temporary data and is often used for function calls and returning function results. It is called a stack because it follows the Last In, First Out (LIFO) principle, meaning that the most recent data added to the stack is the first data to be removed. The stack pointer is used to keep track of the top of the stack and to allow the CPU to push and pop data from the stack as needed. In addition to the program counter and stack pointer, other special registers may include a status register, a memory address register, and a memory data register. These registers are used to control the operation of the CPU and to facilitate communication between the CPU and other components of the computer.

Finally, some CPUs also have a register called the instruction pointer or instruction address register, which contains the address of the current instruction being executed. This is different from the program counter, which points to the next instruction to be fetched. The instruction pointer is used to store the address of a specific instruction in memory, allowing the CPU to jump to that instruction and execute it as part of a branching operation. This is an important feature in programming, allowing for the creation of loops, conditional statements, and other control structures.

Pipelines are a common feature in modern CPUs as they allow for faster processing of instructions by allowing multiple instructions to be executed simultaneously. However, this also means that the underlying complexity of the machine is exposed to compiler and operating system writers, who must take these complexities into account when writing their software. This can be challenging as the behavior of instructions in a pipeline can be difficult to predict, especially when considering the interaction between multiple instructions. As a result, compiler and operating system writers must be careful to ensure that their software takes advantage of the performance benefits offered by pipelines, while also dealing with the challenges they present.

The concept of pipelines and multiple execution units has led to the development of complex instruction set computers (CISC) and reduced instruction set computers (RISC). CISC CPUs have a large and varied instruction set, allowing them to perform a wide range of tasks using fewer instructions. This can be more efficient, as it reduces the number of times the CPU needs to access memory to fetch instructions. RISC

CPUs, on the other hand, have a smaller and more limited instruction set, but make up for this with their ability to process instructions faster through the use of pipelines and multiple execution units.

In addition to these architectural differences, CPUs also differ in their word size, or the number of bits they can process at a time. This determines the maximum size of a memory address and the size of data types that can be stored in registers. A larger word size allows for faster processing of data, but also requires more transistors and consumes more power.

Additionally, modern CPUs often use techniques such as out-of-order execution and speculative execution to further improve performance. Out-of-order execution allows the CPU to execute instructions in a different order than they appear in the program, as long as the final result is not affected. This can be useful if the CPU encounters an instruction that is dependent on the result of a previous instruction that has not yet completed execution. Speculative execution involves the CPU executing instructions based on a prediction of what data will be needed in the future. If the prediction is correct, this can improve performance by allowing the CPU to start executing instructions before all the necessary data is available. However, if the prediction is incorrect, the speculatively executed instructions must be discarded and the correct instructions must be executed instead, which can lead to a decrease in performance.

"A superscalar CPU is able to execute multiple instructions at the same time by utilizing these multiple execution units. This allows for higher levels of parallelism and improved performance compared to a traditional pipeline design. However, this also adds additional complexity for compiler writers and operating system developers as they must carefully consider how to effectively utilize the multiple execution units in order to maximize performance."

One of the key features of an operating system is the ability to switch between kernel mode and user mode. Kernel mode, also known as privileged mode, allows the operating system to have complete control over the hardware and software resources of the computer. This mode is used to execute system-level functions, such as handling interrupts and managing memory. In contrast, user mode is a restricted mode that allows user programs to access the hardware and software resources of the computer, but only through certain predefined interfaces. This separation of privilege helps to protect the operating system and ensure the stability and security of the system. On desktop and server machines, the operating system typically runs in kernel mode, while user programs run in user mode. On embedded systems, a small piece of the operating system may run in kernel mode, while the rest runs in user mode. This distinction between kernel mode and user mode is an important aspect of modern operating systems and helps to ensure the stability, security, and reliability of the system.

The system call is a way for a user program to request services from the operating system. When a user program wants to make a system call, it uses the TRAP instruction to switch from user mode to kernel mode and initiate the operating system. The operating system then performs the requested service and returns control to the user program once it is finished. System calls are an important feature of an operating system, as they allow user programs to access resources and functionality that are only available in the

kernel. This separation of user programs and the kernel helps to ensure the stability and security of the system.

In the world of computers, a trap is an exception or interrupt that occurs when something unexpected happens during the execution of a program. Traps can be caused by a variety of different things, including hardware errors, software bugs, or attempts to access forbidden memory locations. When a trap occurs, the operating system is notified and takes control of the computer in order to handle the exception. This can involve stopping the program that caused the trap, displaying an error message, or taking some other action to prevent further problems. Traps are an important part of the way that computers work, as they allow the operating system to respond to unexpected events and keep the computer running smoothly.

Multithreaded and Multicore Chips